



Photo by **Cris Ovalle** (<https://unsplash.com/@crisovalle>) on **Unsplash** (<https://unsplash.com>)

Code optimisation in Julia

🕒 13 minute read

In this guide we will deal with some key points to write efficient Julia code and I will give you some suggestion on how to identify and deal with code bottlenecks.

Although it is not necessary, I suggest you to run this code inside the Juno IDE. If you don't know what Juno is, I suggest you to take a look at [this](https://techytok.ml/atom-and-juno-setup-for-julia/) (<https://techytok.ml/atom-and-juno-setup-for-julia/>) guide on how to install and use it!

You can find all the code for this guide [here](https://github.com/aurelio-amerio/techytok-examples/tree/master/julia-code-optimization) (<https://github.com/aurelio-amerio/techytok-examples/tree/master/julia-code-optimization>).

Julia is fast but it needs your help!

Julia can be extraordinarily fast, but in order to achieve top speed you need to understand what makes Julia fast.

Contrarily to what may seem at first glance, Julia doesn't draw its speed from **type annotations** but from **type stability**. But what are type annotations and what is type stability?

Type annotations

Type annotation are a way to tell Julia the type that a function should expect and they are useful when you plan to write multiple methods (**multiple dispatch**) for a single function. With type annotations, you can tell a function how to behave according to the type of the arguments that it is given.

For example:

```

1  function test1(x::Int)
2      println("$x is an Int")
3  end
4
5  function test1(x::Float64)
6      println("$x is a Float64")
7  end
8
9  function test1(x)
10     println("$x is neither an Int nor a Float64")
11 end

```

```

1  >>>test1(1)
2  1 is an Int
3
4  >>>test1(1.0)
5  1.0 is a Float64
6
7  >>>test1("techtok")
8  techtok is neither an Int nor a Float64

```

But writing a function like:

```

1  function test2(x::Int, y::Int)
2      return x + y
3  end
4
5  function test2(x::Float64, y::Float64)
6      return x + y
7  end

```

will not yield a performance gain over simply writing `x+y` directly and it is not even a good programming practice in Julia, as it would potentially limit the usage of the function with other types which may be supported indirectly.

Julia is pretty good at doing type inference at run-time and will compile the proper code to handle any type of `x` and `y` or die trying, in the sense that Julia will tell you that it doesn't know how to properly handle the type of `x` and `y`, so it is better to **write code as generic as possible** (i.e. without type annotations) and only use them when multiple dispatch is needed or we know that a function can work *only* with one particular input type.

When you are writing a function which expects a number as an input, it is advisable to use type annotations with **Abstract Types** (<https://docs.julialang.org/en/v1/manual/types/index.html#Abstract-Types-1>), for example using `function test(x::Real)` instead of writing a method for each concrete type. This way the code will be more readable and more generic, a win-win situation! What's more, an user who wants to implement a custom number type, if the type is properly defined, will find that your function will work also for their code!

So if it's not type annotations, which is the first difference which you will spot when comparing for example `c++` to `python`, what makes Julia fast? Roughly said, Julia can compile efficient machine code only if it can infer properly the type of the returned value, which means that your code must be **type stable** if you want to achieve the maximum possible speed.

Type stability

What does type stability means? It means that **the type of the returned value of a function must depend only on the type of the input of the function and not on the particular value it is given.**

Take as an example this function:

```
1  """
2  Returns x if x>0 else returns 0
3  """
4  function test3(x)
5      if x>0
6          return x
7      else
8          return 0
9      end
10 end
```

Let's type the following code in the REPL:

```
1  >>>test3(2)
2  2
3
4  >>>test3(-1)
5  0
```

```
1  >>>test3(2.0)
2  2.0
3
4  >>>test3(-1.0)
5  0
```

Which is a huge problem: as you can see when the input is an `Int` the return output is always and `Int` , but when the input is a `Float64` (like 2.0 or -1.0), the output may either be a `Float64` or an `Int` . In this case Julia cannot determine beforehand whether the returned value will be a `Float64` or an `Int` , which will ultimately lead to a slower code. In this case the problem may be solved by changing line 6 and writing `zero(x)` which returns a “zero” of the same type of `x`.

```

1  """
2  Returns x if x>0 else returns 0
3  """
4  function test4(x)
5      if x>0
6          return x
7      else
8          return zero(x)
9      end
10 end

```

```

1  >>>test4(-1.0)
2  0.00

```

In this case it was “easy” to find the bit of code which leads to type instability, but often the code is more complex and so Julia comes in our help with your new best friend, the `@code_warntype` macro!

What it does is really precious: it tells us if the type of the returned value is unstable and, moreover, it will colour code the result in red if any type instability issue is found. Just call the function with a value which you suspect that may be type unstable and it will do the rest!

```

1  >>>@code_warntype test3(1.0)
2  Body::Union{Float64, Int64} #in red in the REPL
3  1 - %1 = (x > 0)::Bool
4  └─      goto #3 if not %1
5  2 -      return x
6  3 -      return 0
7
8  >>>@code_warntype test4(1.0)
9  Body::Float64 #in blue in the REPL
10 1 - %1 = (x > 0)::Bool
11 └─      goto #3 if not %1
12 2 -      return x
13 3 - %4 = Main.zero(x)::Core.Compiler.Const(0.0, false)
14 └─      return %4

```

As you can see, in the first case the output is either `Float64` or `Int` , which is not good, and in the second case the output can only be a `Float64` (which is good).

In this case Julia may still be able to “recover” from this type instability issue, but if by any chance you see somewhere an `::Any` (which will always be marked in red) you will know that there is some serious type instability issue which needs to be fixed. Sometimes fixing type instability is easy, sometimes it is not, you will have to deal with it case by case, but the reward is always great, in my programming experience I have seen a **performance gain of a thousand times** by fixing type stability issues in my code!

Another common error is changing the type of a variable inside a function:

```

1  function test5()
2      r=0
3      for i in 1:10
4          r+=sin(i)
5      end
6      return r
7  end
8
9  >>>@code_warntype test5()
10 Variables
11  #self#::Core.Compiler.Const(test5, false)
12  r::Union{Float64, Int64}
13  @_3::Union{Nothing, Tuple{Int64,Int64}}
14  i::Int64
15
16  Body::Float64
17  1 -      (r = 0)
18  |   %2  = (1:10)::Core.Compiler.Const(1:10, false)
19  |       (@_3 = Base.iterate(%2))
20  |   %4  = (@_3::Core.Compiler.Const((1, 1), false) === nothing)::Core.Compiler.Const(false,
21  false)
22  |   %5  = Base.not_int(%4)::Core.Compiler.Const(true, false)
23  |   └─   goto #4 if not %5
24  |       (i = Core.getfield(%7, 1))
25  |   %9  = Core.getfield(%7, 2)::Int64
26  |   %10 = r::Union{Float64, Int64}
27  |   %11 = Main.sin(i)::Float64
28  |       (r = %10 + %11)
29  |       (@_3 = Base.iterate(%2, %9))
30  |   %14 = (@_3 === nothing)::Bool
31  |   %15 = Base.not_int(%14)::Bool
32  |   └─   goto #4 if not %15
33  3 -      goto #2
34  4 ---      return r::Float64

```

As you can see `r` at the beginning is an `Int64` and then is converted into a `Float64`, which slows down the function. In this case it is enough to declare `r=0.00` to solve the problem.

Function profiling

In this section we will learn how to find the bottlenecks in the execution of a function, so that we know which parts of the function should be optimised.

As you probably already know, it is possible to measure the execution time of a function using the `@time` macro or preferably `@btime` (which is included in the package `BenchmarkTools`). They are useful when we want to measure a single function call, but they give no information on what is making a function slow. For this reason we need a tool that enables us to identify which line of code is responsible for the bottleneck.

Luckily there are two exceptional packages which help us in profiling: `Profile` and `Juno` with the **Juno IDE**. Both of the tools will run the desired function once and will log the execution time of each line of code.

`Profile` works completely in the REPL and will produce a log file, while `Juno` works only inside the Juno IDE, and will display some information in a graph.

Profile

Let's write two new functions which perform heavy calculations:

```
1  function take_a_breath()
2      sleep(22*1e-3)
3      return
4  end
5
6  function test8()
7      r=zeros(100,100)
8      take_a_breath()
9      for i in 1:100
10         A=rand(100,100)
11         r+=A
12     end
13     return r
14 end
```

We shall now profile `test8` :

```
1  using Profile
2  test8()
3  Profile.clear()
4  @profile test8()
5  Profile.print()
```

If you see a log file extraordinarily long, run line 2 to 4 again as it is possible that you have profiled the compilation of some functions and not `test8`.

You should see something like:

```

8 ...amples\julia-code-optimization\code-optimization.jl:101; test8()
4 ...d\usr\share\julia\stdlib\v1.2\Random\src\Random.jl:265; rand
4 ...usr\share\julia\stdlib\v1.2\Random\src\Random.jl:244; rand!
4 ...ld\usr\share\julia\stdlib\v1.2\Random\src\RNGs.jl:456; rand!
4 ...ld\usr\share\julia\stdlib\v1.2\Random\src\RNGs.jl:450; _rand!
4 .\gcutils.jl:87; macro expansion
4 ...d\usr\share\julia\stdlib\v1.2\Random\src\RNGs.jl:438; rand! (::Random.MersenneTwister, :
4 ...usr\share\julia\stdlib\v1.2\Random\src\RNGs.jl:412; fill_array!
4 ...usr\share\julia\stdlib\v1.2\Random\src\DSFMT.jl:95; dsfmt_fill_array_close_open! (::Ran
1 ...amples\julia-code-optimization\code-optimization.jl:104; test8()
1 .\arraymath.jl:47; + (::Array{Float64,2}, ::Array{Float64,2})
1 .\broadcast.jl:752; broadcast (::typeof(+), ::Array{Float64,2}, ::Array{Float64,2})
1 .\boot.jl:406; materialize

```

The number 8 (which may vary) is an indication of how long a function runs. We also have a 1 and a 4, so probably this block is the one responsible for the bottleneck. Profile thus gives you a hint that at line 101 (line 8 of the code snippet) there is a bottleneck: who would have imagined that sleeping for 22ms would have been a slow down?

Unfortunately using `Profile.print()` is not really user friendly and for this reason Juno comes in our help!

Before trying the next steps please restart the Julia REPL as I have found that sometimes the Juno profiler doesn't work properly if you have already imported `Profile`.

After having defined `test8` again, call the Juno `@profiler`

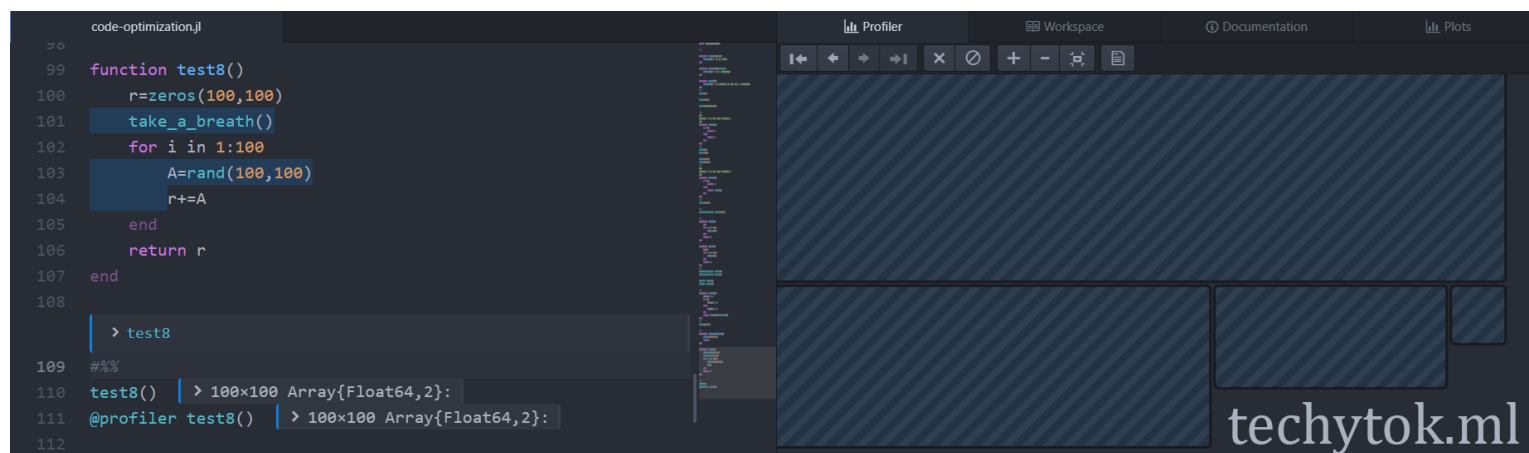
```

1 test8()
2 @profiler test8()

```

Again, you may need to call `@profiler test8()` a few times before you get the right results which don't include the compilation process.

If everything goes right, you should see something like this:



As you can see there is a new panel on the right called “Profiler” which shows three smaller blocks: you can click on them and it will lead to the piece of code responsible for the relative call: the bigger the block, the longer it takes to run the piece of code. There will also be some bars on top of your code: they show which

line of code has a greater impact on the run time (as you can see on the left in the picture), which is really useful to find at a glance where is the bottleneck!

Tips and Tricks

Now that we have found how to identify type instability issues and how to profile your code, I will give you some tips on how you can make your code faster by changing some little details.

@inbounds

When you are working with **for loops** and you are sure that the loop condition will not lead to out of bound errors, you can use the `@inbounds` macro to skip bound checking and gain a speed boost

```
1  using BenchmarkTools
2  arr1=zeros(10000)
3
4  function put1!(arr)
5      for i in 1:length(arr)
6          arr[i]=1.0
7      end
8  end
9
10 function put1_inbounds!(arr)
11     @inbounds for i in 1:length(arr)
12         arr[i]=1.0
13     end
14 end
15
16 >>>@btime put1!($arr1)
17 3.814 μs (0 allocations: 0 bytes)
18
19 >>>@btime put1_inbounds!($arr1)
20 1.210 μs (0 allocations: 0 bytes)
```

On my PC `put1!` takes 3.814μs, while `put1_inbounds!` takes only 1.220μs (which is a noticeable difference).

Static Arrays

When the dimension of an array is known beforehand and it will not change over time, static arrays become incredibly powerful as they enable the compiler to perform advanced optimisations.

They are particularly useful when you need to perform operations (like the sum of the members of an array or matrix operations) on small vectors (<100 elements) or matrices.

If we run the micro benchmark from the official [github repository](https://github.com/JuliaArrays/StaticArrays.jl).

(https://github.com/JuliaArrays/StaticArrays.jl/blob/master/perf/README_benchmarks.jl), we get something like this:

```
1  =====
2      Benchmarks for 3x3 Float64 matrices
3  =====
4  Matrix multiplication          -> 7.4x speedup
5  Matrix multiplication (mutating) -> 2.4x speedup
6  Matrix addition              -> 31.5x speedup
7  Matrix addition (mutating)   -> 2.5x speedup
8  Matrix determinant           -> 121.2x speedup
9  Matrix inverse               -> 70.3x speedup
10 Matrix symmetric eigendecomposition -> 22.6x speedup
11 Matrix Cholesky decomposition -> 7.2x speedup
12 Matrix LU decomposition       -> 8.4x speedup
13 Matrix QR decomposition       -> 47.1x speedup
14  =====
```

As you can see the performance gain is enormous!

Matrix indices

When writing performant code, we need to keep in mind how data is stored in the memory. In particular, while in c++ and python matrices are stored in memory in a “row first” manner, in Julia they are stored column first. Which means that when we cycle the indices of a matrix in a nested for loop, if we want to read or write on the matrix, it is better to do so “by column” instead of “by row”.

Here is an example:

```
1  arr1 = zeros(100, 200)
2
3  @btime for i in 1:100
4      for j in 1:200
5          arr1[i,j] = 1
6      end
7  end
8  # 386.600 μs
9
10 @btime for j in 1:200
11     for i in 1:100
12         arr1[i,j] = 1
13     end
14 end
15 # 380.800 μs
```

Keyword arguments

When a function is in a critical inner loop, avoid keyword arguments and use only positional arguments, this will lead to better optimisation and faster execution.

```
1 | function foo1(x,y=2) ... # preferable
2 |
3 | function foo2(x; y=2) ... # discouraged inside inner loops
```

For example, if we define the following functions:

```
1 | function test_positional(a, b, c)
2 |     return a + b + c
3 | end
4 |
5 | function test_keyword(;a, b, c)
6 |     return a + b + c
7 | end
```

We can see the LLVM code which gets compiled for each one using the `@code_llvm` macro.

```
1 | >>>@code_llvm test_positional(1,2,3)
2 |
3 | ; @ D:\Users\Aure\Documents\GitHub\techtok-examples\julia-code-optimization\code-
4 | optimization.jl:210 within `test_positional'
5 | ; Function Attrs: uwtable
6 | define i64 @julia_test_positional_18504(i64, i64, i64) #0 {
7 |   top:
8 |   ;  @ operators.jl:529 within `+' @ int.jl:53
9 |   %3 = add i64 %1, %0
10 |   %4 = add i64 %3, %2
11 |   ; L
12 |   ret i64 %4
}
```

```

1  @code_llvm test_keyword(a=1,b=2,c=3)
2
3  ; @ none within `#test_keyword'
4  ; Function Attrs: uwtable
5  define i64 @"julia_#test_keyword_18505"({ i64, i64, i64 } addrspace(11)* nocapture nonnull
6  readonly dereferenceable(24)) #0 {
7  top:
8  ;  ⌈ @ namedtuple.jl:107 within `getindex'
9      %1 = getelementptr inbounds { i64, i64, i64 }, { i64, i64, i64 } addrspace(11)* %0, i64 0, i32
10  2
11      %2 = getelementptr inbounds { i64, i64, i64 }, { i64, i64, i64 } addrspace(11)* %0, i64 0, i32
12  1
13      %3 = getelementptr inbounds { i64, i64, i64 }, { i64, i64, i64 } addrspace(11)* %0, i64 0, i32
14  0
15  ;  ⌋
16  ;  ⌈ @ D:\Users\Aure\Documents\GitHub\techtok-examples\julia-code-optimization\code-
17  optimization.jl:214 within `#test_keyword#9'
18  ;  ⌈⌈ @ operators.jl:529 within `+' @ int.jl:53
19      %4 = load i64, i64 addrspace(11)* %3, align 8
20      %5 = load i64, i64 addrspace(11)* %2, align 8
21      %6 = add i64 %5, %4
22      %7 = load i64, i64 addrspace(11)* %1, align 8
23      %8 = add i64 %6, %7
24  ;  ⌋⌋
25      ret i64 %8
26  }

```

As you can see, when we call `test_keyword` there are 9 more lines of code. If `test_keyword` is a function inside a critical inner loop this will lead to a considerable slowdown.

Avoid global scope variables

As the title says, try to avoid global scope variables as much as you can, as it is more difficult for Julia to optimise them. If possible, try to declare them as `const` : most of the times a global variable is likely to be a constant, which is a good solution to in part alleviate the problem. When possible, pass all the required data to the function by argument, this way the function will be more flexible and better optimised.

Conclusion

These are some tips and tricks to make your Julia code run faster! The **take home message** for good performance is **keep your code as general as possible and as simple as possible** and focus on **type stability**.

Remember: speed in Julia comes from **type stability** and **not type annotations**, which means that you should not use type annotations if not required and that the **returned value must depend only on the type of the arguments**.

If you liked this lesson and you would like to receive further updates on what is being published on this website, I encourage you to subscribe to the **newsletter** (<https://techytok.com/newsletter/>)! If you have any **question** or **suggestion**, please post them in the **discussion below**!

Thank you for reading this lesson and see you soon on TechyTok!

Tags: Julia Optimisation

Categories: Guide Tutorial

Updated: October 26, 2019

LEAVE A COMMENT

ALSO ON TECHYTOK

Variable Scope

7 months ago • 4 comments

From zero to Julia Lesson 6. Variable scope

Control Flow

8 months ago • 5 comments

From zero to Julia Lesson 4. Control Flow

Multiprocessing i Julia: writing a m

9 months ago • 3 comm

A tutorial on how to w module in Julia which multiprocessing

7 Comments techytok Privacy Policy 1 Login Recommend 3 Tweet Share Sort by Best

Join the discussion...

LOG IN WITH

OR SIGN UP WITH DISQUS

Name

Dictino • 6 months ago

Hi Aurelio

Very nice posts, I learn a lot of things reading your lessons, thanks for your effort.

I want to point out that there is an important optimization missing ie. "put things into functions".

For instance in your @inbounds example the benefit is not very significant because you are working at global scope if you define functions there is a

huge difference:

using BenchmarkTools

```
arr1=zeros(10000)
```

```
@btime for i in 1:10000
```

```
arr1[i] = 1
```

```
end
```

```
# 300.028 us (0480 allocations: 148.27 KiB)
```

[see more](#)

^ | v • Reply • Share ›



aureamerio Mod ➔ Dictino • 6 months ago

You are right! Thank you, I have updated the lesson!

Just a reminder: when you use `@btime` always interpolate the variables, i.e. call `@btime put_one!($arr1)` and not `put_one!(arr1)`, as the results may not be accurate otherwise (in my case, interpolating the variable I obtained a 3x speedup between the function with and without `@inbounds`).

^ | v • Reply • Share ›



Dictino ➔ aureamerio • 6 months ago

My bad, I always forget the interpolation :(

Anyway It's very nice to see that the function is even faster that I've posted :)

^ | v • Reply • Share ›



purjus • 6 months ago

Great post! I just have one question. Why are keyword arguments costly? Can you expand on this?

^ | v • Reply • Share ›



aureamerio Mod ➔ purjus • 6 months ago • edited

It is not much, but the generated llvm code for a function with keyword arguments is longer than the exact same function with only positional arguments.

Take for example the following functions:

```
function test_positional(a, b, c)
    return a + b + c
end
```

```
function test_keyword(;a, b, c)
    return a + b + c
end
```

If you type `@code_llvm test_positional(1,2,3)` you get:

```
; @ D:\Users\Aure\Documents\GitHub\techtok-examples\julia-
code-optimization\code-optimization.jl:210 within `test_positional'
```

```
| ; Function Attrs: uwtable  
| define i64 @iulia_test_positional_18486(i64 i64 i64) #0 {
```

[see more](#)

2 ^ | v • Reply • Share ›



D ZJ • 9 months ago

This is great. Something that is unintuitive about Julia is that sometimes it's better not to encode type info in your type to help with type inference time. This is so that the compiler doesn't spend too much specialising code to all the types in your object. E.g. DataFrames.jl vs IndexedTable.jl. Perhsp talk abt that

^ | v • Reply • Share ›



aureamerio Mod ➔ **D ZJ** • 9 months ago

You are right, I don't know much about indexed tables, but I will definitely take a look at that and write an article about type annotations in the future!